

# Reading and Writing Arithmetic: Automating Ring Equalities in Agda

ANONYMOUS AUTHOR(S)

We present a new library which automates the construction of equivalence proofs between polynomials over commutative rings and semirings in the programming language Agda [Norell and Chapman 2008]. It is asymptotically faster than Agda's existing solver. We use reflection to provide a simple interface to the solver, and demonstrate a novel use of the constructed relations: step-by-step solutions.

Additional Key Words and Phrases: proof automation, equivalence, proof by reflection, step-by-step solutions

`lemma` :  $\forall x y \rightarrow x + y * 1 + 3 \approx 2 + 1 + y + x$

`lemma`  $x y =$  `begin`

$x + y * 1 + 3 \approx \langle \text{refl } \langle +\text{-cong} \rangle * \text{-identity}^r y \langle +\text{-cong} \rangle \text{refl } \{3\} \rangle$

$x + y + 3 \approx \langle +\text{-comm } x y \langle +\text{-cong} \rangle \text{refl} \rangle$

$y + x + 3 \approx \langle +\text{-comm } (y + x) 3 \rangle$

$3 + (y + x) \approx \langle \text{sym } (+\text{-assoc } 3 y x) \rangle$

$2 + 1 + y + x \blacksquare$

`lemma` = `solve` `NatRing`

(a) A Tedious Proof

(b) The Solver

Fig. 1. Comparison Between A Manual Proof and The Automated Solver

## 1 INTRODUCTION

Doing mathematics in dependently-typed programming languages like Agda has a reputation for being tedious, awkward, and difficult. Even simple arithmetic identities like the one in Fig. 1 require fussy proofs (Fig. 1a).

This need not be the case! With some carefully-designed tools, mathematics in Agda can be easy, friendly, and fun. This work describes one such tool: an Agda library which automates the construction of these kinds of proofs, making them as easy as Fig. 1b.

As you might expect, our solver comes accompanied by a formal proof of correctness. Beyond that, though, we also strove to satisfy the following requirements:

**Friendliness and Ease of Use** Proofs like the one in Fig. 1a aren't just boring; they're *difficult*.

The programmer needs to remember the particular syntax for each step ("is it `+comm` or `+commutative?`"), and often they have to put up with poor error messages.

We believe this kind of difficulty is why Agda's current ring solver [Danielsson 2018] enjoys little widespread use. Its interface (Fig. 2) is almost as verbose as the manual proof, and it requires programmers to learn another syntax specific to the solver.

Our solver strives to be as easy to use as possible: the high-level interface is simple (Fig. 1b), we don't require anything of the user other than an implementation of one of the supported algebras, and effort is made to generate useful error messages.

```
lemma = +-*-Solver.solve 2 (λ x y → x:+ y:* con 1 :+ con 3 := con 2 :+ con 1 :+ y:+ x) refl
```

Fig. 2. The Old Solver

**Performance** Typechecking dependently-typed code is a costly task. Automated solvers like the one presented here can greatly exacerbate this cost: in our experience, it wasn't uncommon for Agda's current ring solver to spend upwards of 10 minutes proving a single identity.

In practice, this means two things: firstly, large libraries for formalizing mathematics (like Meshveliani [2018]) can potentially take hours to typecheck (by which time the programmer has understandably begun to reconsider the whole notion of mathematics on a computer); secondly, certain identities can simply take too long to typecheck, effectively making them "unprovable" in Agda altogether!

The kind of solver we provide here is based on Coq's [Team 2018] ring tactic, described in Grégoire and Mahboubi [2005]. While we were able to apply the same optimizations that were applied in that paper, we found that the most significant performance improvements came from a different, and somewhat surprising part of the program. The end result is that our solver is asymptotically (and practically) faster than Agda's current solver.

**Educational Features** While our solver comes with the benefit of formal correctness, it's still playing catch-up to other less-rigorous computer algebra systems in terms of features. These features have driven systems like Wolfram|Alpha [Wolfram Research, Inc. 2019] to widespread popularity among (for instance) students learning mathematics.

We will take just one of those features ("pedagogical", or step-by-step solutions [The Development Team 2009]), and reimplement it in Agda using our solver. In doing so, we will explore some of the theory behind it, and present a formalism that describes the nature of "step-by-step" solutions.

## 2 OVERVIEW OF THE PROOF TECHNIQUE

There are a number of ways we can automate proofs in a dependently-typed programming language, including Prolog-like proof search [Kokke and Swierstra 2015], Cooper's algorithm over Presburger arithmetic [Allais 2011], etc. Here, we will use a reflexive technique [Boutin 1997] in combination with sparse Horner Normal Form. The high-level diagram of the proof strategy is presented in Fig. 3.

The identity we'll be working with is the lemma in Fig. 1: the left and right hand side of the equality are at the bottom of the diagram. Our objective is to link those two expressions up through repeated application of the ring axioms. We do this by converting both expressions to a normal form (seen at the top of the diagram), and then providing a proof that this conversion is correct according to the ring axioms (the `correct` function in the diagram). Finally, we link up all of these proofs, and if the two normal forms are definitionally equal, the entire thing will typecheck, and we will have proven the equality.

### 2.1 The Expr AST

In Agda, we can't manipulate the expressions we want to prove directly: instead, we will construct an AST for each expression, and then do our manipulation on that.

The AST type (`Expr`) has a constructor for each of the ring operators, as well as constructors for both variables and constants. The ASTs for both expressions we want to prove can be seen on either side of Fig. 3. Constants are constructed with `K`, and variables are referred to by their de Bruijn index (so `x` becomes `! 0`).




Fig. 3. The Reflexive Proof Process

From here, we can “evaluate” the AST in one of two ways: in a non-normalized way ( $\llbracket \_ \rrbracket$ ), or in a normalizing way ( $\llbracket \_ \rrbracket$ ). This means that the goal of the **correct** function is to show equivalence between  $\llbracket \_ \rrbracket$  and  $\llbracket \_ \rrbracket$ .

Finally, we *don't* want to force users to construct the **Expr** AST themselves. This is where reflection comes in: it automates this construction (the path labeled **quoteTerm** in the diagram) from the goal type.

## 2.2 Almost Rings

So far, we have been intentionally vague about the precise algebra we're using. As in Grégoire and Mahboubi [2005, section 5], we use an algebra called an *almost-ring*. It has the regular operations (+, \* (multiplication), -, 0, and 1), such that the following equations hold:

$$0 + x = x \quad (1)$$

$$x + y = y + x \quad (2)$$

$$x + (y + z) = (x + y) + z \quad (3)$$

$$1 * x = x \quad (4)$$

$$x * y = y * x \quad (5)$$

$$x * (y * z) = (x * y) * z \quad (6)$$

$$(x + y) * z = x * z + y * z \quad (7)$$

$$0 * x = 0 \quad (8)$$

$$-(x * y) = -x * y \quad (9)$$

$$-(x + y) = -x + -y \quad (10)$$

The equations up to 8 represent a pretty standard definition of a (commutative) semiring. To get from there to a ring, we would usually ask for something like:

$$x + -x = 0 \quad (11)$$

So why do we instead ask for 9 and 10?

The reason is *flexibility*. Under this formulation, we can admit types like  $\mathbb{N}$  which don't have additive inverses. Instead, these types can simply supply the identity function for  $-$ , and then 9 and 10 will still hold.

A potential worry is that because we don't require  $x + -x = 0$  axiomatically, it won't be provable in our system. Happily, this is not the case: as long as  $1 + -1$  reduces to 0 in the coefficient set, the solver will verify the identity.

In the library, the algebra is represented by the `AlmostCommutativeRing` type, a record with fields for each of the ring axioms, defined over a user-supplied equivalence relation. Just as in Agda's current solver, we also ask for one extra function: a weakly decidable predicate to test if a constant is equal to zero.

```
is-zero : ∀ x → Maybe (0# ≈ x)
```

This function is used to speed up some internal algorithms in the solver, but it isn't an essential component. By making it *weakly* decidable, we allow users to skip it (`is-zero = const nothing`) if their type doesn't support decidable equality, or provide it (and get the speedup) if it does.

### 3 THE INTERFACE

A decent interface is crucial if we want the solver to be broadly useful. We strove to make our interface as simple and *small* as possible. Aside from the `AlmostCommutativeRing` type described above, the user-facing portion of our library consists of just two macros: `solve` and `solveOver`. Each of these infer the goal from their context, and automatically construct the required machinery to prove the equality.

`solve` is demonstrated in Fig. 1b. It takes a single argument: an implementation of the algebra. `solveOver` is designed to be used in conjunction with manual proofs, so that a programmer can automate a “boring” section of a larger more complex proof. It is called like so (line 5):

```
1 lemma : ∀ x y → x + y * 1 + 3 ≈ 2 + 1 + y + x
2 lemma x y =
3   begin
4     x + y * 1 + 3 ≈ ⟨ +-comm (x + y * 1) 3 ⟩
5     3 + (x + y * 1) ≈ ⟨ solveOver (x :: y :: []) Nat.ring ⟩
6     3 + y + x ≡ ⟨ ⟩
7     2 + 1 + y + x ■
```

As well as the `AlmostCommutativeRing` implementation, this macro takes a list of free variables to use to compute the solution.

Because this interface is quite simple, it's worth pointing out what we *don't* require from the user:

- We don't ask the user to construct the `Expr` AST which represents their proof obligation. Compare this to Fig. 2: we had to write the type of the proof twice (once in the signature and again in the AST), and we had to learn the syntax for the solver's AST. As well as being more verbose, this approach is less composable: every change to the proof type has to be accompanied by a corresponding change in the call to the solver. In contrast, the call to `solveOver` above effectively amounts to a demand for the compiler to “figure it out!” Any change to the expressions on either side will result in an *automatic* change to the proof constructed.
- We don't ask the user to write any kind of “reflection logic” for their type. In other words, we don't require a function which (for instance) recognizes and parses the user's type in the reflected AST, or a function which does the opposite, converting a concrete value into the AST that (when unquoted) would produce an expression equivalent to the quoted value.




Fig. 4. Time (in seconds) to prove each expression is equal to its expanded form ( $n = 5$  for each).

This kind of logic is complex, and very difficult to get right. While some libraries can assist with the task [Norell 2018; van der Walt and Swierstra 2013] it is still not fully automatic.

Finally, despite the simplicity and ease-of-use described above, the solver is *not* specialized to a small number of types like `N` and so on. The whole library, including the reflection-based interface, will work with any type with an `AlmostCommutativeRing` instance.

## 4 PERFORMANCE

## 5 PEDAGOGICAL PROOFS

## 6 RELATED WORK

## REFERENCES

- G Allais. 2011. Deciding Presburger Arithmetic Using Reflection. (May 2011). <https://gallais.github.io/pdf/presburger10.pdf>
- Samuel Boutin. 1997. Using Reflection to Build Efficient and Certified Decision Procedures. In *Theoretical Aspects of Computer Software (Lecture Notes in Computer Science)*, Martín Abadi and Takayasu Ito (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 515–529.
- Nils Anders Danielsson. 2018. The Agda Standard Library. <https://agda.github.io/agda-stdlib/README.html>
- Benjamin Grégoire and Assia Mahboubi. 2005. Proving Equalities in a Commutative Ring Done Right in Coq. In *Theorem Proving in Higher Order Logics (Lecture Notes in Computer Science)*, Vol. 3603. Springer Berlin Heidelberg, Berlin, Heidelberg, 98–113. [https://doi.org/10.1007/11541868\\_7](https://doi.org/10.1007/11541868_7)
- Pepijn Kokke and Wouter Swierstra. 2015. Auto in Agda. In *Mathematics of Program Construction (Lecture Notes in Computer Science)*, Ralf Hinze and Janis Voigtländer (Eds.). Springer International Publishing, 276–301. <http://www.staff.science.uu.nl/~swier004/publications/2015-mpc.pdf><http://www.staff.science.uu.nl/~swier004/publications/2015-mpc.pdf>
- Sergei D. Meshveliani. 2018. DoCon-A a Provable Algebraic Domain Constructor. <http://www.botik.ru/pub/local/Mechveliani/docon-A/2.02/>
- Ulf Norell. 2018. Agda-Prelude: Programming Library for Agda. <https://github.com/UlfNorell/agda-prelude>
- Ulf Norell and James Chapman. 2008. Dependently Typed Programming in Agda. (2008), 41.
- The Coq Development Team. 2018. The Coq Proof Assistant, Version 8.8.0. <https://doi.org/10.5281/zenodo.1219885>
- The Development Team. 2009. Step-by-Step Math. <http://blog.wolframalpha.com/2009/12/01/step-by-step-math/>
- Paul van der Walt and Wouter Swierstra. 2013. Engineering Proof by Reflection in Agda. In *Implementation and Application of Functional Languages*, Ralf Hinze (Ed.). Vol. 8241. Springer Berlin Heidelberg, Berlin, Heidelberg, 157–173. [https://doi.org/10.1007/978-3-642-41582-1\\_10](https://doi.org/10.1007/978-3-642-41582-1_10)
- Wolfram Research, Inc. 2019. Wolfram|Alpha. Wolfram Research, Inc.. <https://www.wolframalpha.com/>

## A APPENDIX

Text of appendix ...